

Modern Optimization Algorithms in Deep Learning: A Comprehensive Comparative Analysis

Faculty of Artificial Intelligence — Menoufia University

Data Science Fundamentals | 2025–2026

Abstract

Optimization lies at the core of training modern machine learning models. This paper provides a complete, self-contained guide to modern optimization algorithms, progressing from the limitations of basic **Gradient Descent** through classical improvements to state-of-the-art adaptive methods: **Adam**, **AdamW**, **AdaGrad**, **RMSProp**, **AdaDelta**, **Adamax**, **Adafactor**, and **AdaBelief**. Every algorithm is presented with its mathematical formulation, intuitive explanation, historical motivation, strengths, limitations, and practical use cases. A detailed comparative analysis, experimental evaluation on MNIST and CIFAR-10, and in-depth discussion of trade-offs — speed vs. generalization, memory efficiency, numerical stability — are included. **Key insight:** no single optimizer is universally best. The optimal choice depends on dataset size, model architecture, and priorities: fast convergence vs. strong generalization vs. memory constraints.

1. Introduction

1.1 Definition of Optimization in Machine Learning

Optimization in machine learning is the process of iteratively adjusting a model's parameters θ — weights and biases — to minimize a **loss function** $\mathcal{L}(\theta)$. The loss quantifies how far the model's predictions deviate from true labels. The ultimate goal is not only low training loss but also strong **generalization** to unseen data — a distinction that critically influences optimizer choice.

1.2 Role of Gradient Descent

Gradient Descent (GD) is the foundational algorithm for training neural networks. It uses the gradient $\nabla \mathcal{L}(\theta)$ — the direction of steepest increase in loss — to move parameters in the opposite direction, reducing the loss step by step. Virtually every modern deep learning model, from CNNs to Transformers to GANs, is trained using some variant of gradient descent.

1.3 Limitations of Vanilla Gradient Descent

Full-batch (vanilla) gradient descent suffers from three fundamental problems at scale:

- **Slow convergence:** In high-dimensional parameter spaces or elongated loss landscapes (ravines), gradient updates zig-zag inefficiently, making progress painfully slow toward the minimum.
- **Learning rate sensitivity:** A learning rate η that is too small causes glacially slow training; one that is too large causes divergence or oscillation around the minimum.
- **Saddle points and local minima:** Real loss landscapes of deep networks are highly non-convex. Gradients can vanish near saddle points — regions where the gradient is near zero but the point is not a local minimum — trapping optimization for many steps.

These problems become severe in deep networks with millions of parameters trained on noisy, stochastic mini-batches. The research community responded with a sequence of increasingly sophisticated algorithms that address each limitation.

1.4 Motivation for Advanced Optimizers

Momentum was introduced to accelerate progress in consistent gradient directions and dampen oscillations. **Per-parameter adaptive learning rates** (AdaGrad, RMSProp) address the fact that different parameters require very different step sizes. **Bias correction** (Adam) fixes initialization effects that distort adaptive methods in early training. **Decoupled regularization** (AdamW) restores the proper effect of weight decay. **Memory-efficient factorization** (Adafactor) enables training massive Transformer models that would otherwise exceed GPU memory.

2. Mathematical Foundations

2.1 Gradient Descent Update Rule

The full-batch gradient descent update at step t is:

$$\theta_{t+1} = \theta_t - \eta \cdot \nabla \mathcal{L}(\theta_t)$$

where θ is the vector of all model parameters, $\eta > 0$ is the learning rate (step size), and $\nabla \mathcal{L}(\theta_t)$ is the gradient of the loss function evaluated at the current parameters.

2.2 Stochastic and Mini-batch Gradient Descent

Full-batch GD is computationally impractical for large datasets. Stochastic Gradient Descent (SGD) approximates the true gradient using a single randomly selected example:

$$\theta_{t+1} = \theta_t - \eta \cdot \nabla L(\theta_t; x_i, y_i)$$

Mini-batch SGD — using batch sizes of 32–256 — balances the noise of SGD with computational efficiency and is the **default in practice**. The gradient noise introduced by mini-batches acts as implicit regularization, often helping the optimizer escape sharp minima and find flatter, more generalizable solutions.




Figure 1. Loss landscape contours (elongated bowl simulating ill-conditioning). Left: Vanilla GD zigzags inefficiently across the ravine, making slow progress toward the minimum (★). Right: SGD with momentum smoothly follows the valley, converging far more quickly.

3. Classical Optimization Algorithms

3.1 SGD with Momentum

Momentum, introduced by Polyak (1964) and popularized in neural networks by Rumelhart et al. (1986), adds a velocity term v_t that accumulates past gradients:

$$v_t = \beta \cdot v_{t-1} + (1 - \beta) \cdot \nabla L(\theta_t) \quad (\beta \approx 0.9)$$

$$\theta_{t+1} = \theta_t - \eta \cdot v_t$$

The velocity v_t acts like a ball rolling down a hill: it builds speed in consistent directions and resists changes in direction. This dramatically reduces oscillation in ravine-like loss landscapes and accelerates convergence in directions of persistent curvature.

3.2 RMSProp

Introduced by Tieleman & Hinton (2012), RMSProp solves AdaGrad's aggressive learning rate decay by replacing the cumulative sum with an **exponentially decaying moving average** of squared gradients:

$$s_t = \beta \cdot s_{t-1} + (1 - \beta) \cdot (\nabla L(\theta_t))^2 \quad (\beta \approx 0.9)$$

$$\theta_{t+1} = \theta_t - (\eta / \sqrt{s_t + \epsilon}) \cdot \nabla L(\theta_t) \quad (\epsilon = 1e-8)$$

Parameters with consistently large gradients receive smaller updates; those with small gradients receive larger ones. This **per-parameter adaptation** makes RMSProp particularly effective for recurrent neural networks (RNNs), where gradients can vary wildly across time steps.

4. Modern Optimization Algorithms

4.1 Adam — Kingma & Ba (2015)

Adam (Adaptive Moment Estimation) is the most widely used optimizer in deep learning. It combines momentum (1st moment) with RMSProp-style gradient scaling (2nd moment) and adds **bias correction** to account for the zero-initialization of both moment estimates:

$$m_t = \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot \nabla L(\theta_t) \quad [1st\ moment]$$

$$v_t = \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot (\nabla L(\theta_t))^2 \quad [2nd\ moment]$$

$$\hat{m}_t = m_t / (1 - \beta_1^t) \quad [bias\ correction]$$

$$\hat{v}_t = v_t / (1 - \beta_2^t) \quad [bias\ correction]$$

$$\theta_{t+1} = \theta_t - \eta \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \varepsilon) \quad [update]$$

Default hyperparameters $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\varepsilon = 1e-8$ work well for most tasks with minimal tuning. Adam typically converges 2–5× faster than SGD on deep learning benchmarks.




Figure 2. Adam internal dynamics over 60 training steps. Left: Raw gradient (dashed) vs. 1st moment m_t (solid) and bias-corrected $\hat{m}_t$ (dotted) — momentum smooths noisy gradients. Center: Gradient² and 2nd moment v_t track gradient variance. Right: The effective per-step learning rate is automatically scaled by gradient magnitude, staying well below the global η .

4.2 AdamW — Loshchilov & Hutter (2019)

A critical flaw of Adam is how it handles **L2 regularization (weight decay)**. When weight decay is added to the loss function in the standard way, it enters the gradient and is *scaled by the adaptive denominator* $1/\sqrt{\hat{v}_t}$. This distorts the regularization effect: parameters that update frequently (large $\hat{v}_t$) receive far less regularization than intended.

AdamW fixes this by **decoupling weight decay from the gradient update** and applying it directly to the parameters:

$$\theta_{t+1} = \theta_t - \eta \cdot (\hat{m}_t / (\sqrt{\hat{v}_t} + \varepsilon)) + \lambda \cdot \theta_t$$

where λ is the weight decay coefficient applied uniformly to all parameters, regardless of gradient history. This restores proper **L2 regularization** and dramatically improves generalization in Transformer architectures. AdamW is now the **de facto default optimizer** for large language models, vision transformers, and most modern deep learning.

4.3 AdaGrad — Duchi, Hazan & Singer (2011)

AdaGrad (Adaptive Gradient) was the first optimizer to introduce per-parameter adaptive learning rates. It accumulates the sum of all squared gradients from the beginning of training:

$$G_t = \sum_{i=1}^t (\nabla L(\theta_i))^2$$
$$\theta_{t+1} = \theta_t - (\eta / \sqrt{G_t + \epsilon}) \cdot \nabla L(\theta_t)$$

This is highly effective for **sparse features** (e.g., rare words in NLP): infrequently updated parameters accumulate little G_t , so they receive large effective learning rates; frequent parameters converge quickly. The critical limitation is that G_t grows monotonically, causing the effective learning rate to **decay to zero**, permanently halting learning in long training runs.

4.4 AdaDelta — Zeiler (2012)

AdaDelta extends AdaGrad by replacing the cumulative sum G_t with a **sliding window** of recent squared gradients, implemented as an exponentially decaying average:

$$E[g^2]_t = \rho \cdot E[g^2]_{t-1} + (1 - \rho) \cdot g_t^2$$
$$\Delta\theta_t = -(\sqrt{E[\Delta\theta^2]_{t-1} + \epsilon} / \sqrt{E[g^2]_t + \epsilon}) \cdot g_t$$

A key innovation is that AdaDelta uses the **RMS of previous parameter updates** $E[\Delta\theta^2]$ as the numerator, eliminating the need for a manually specified global learning rate η entirely. In practice, however, AdaDelta often converges more slowly than Adam.

4.5 Adamax — Kingma & Ba (2015)

Adamax is a variant of Adam that replaces the L2-norm scaling of the 2nd moment with the **L^∞ (infinity) norm**:

$$u_t = \max(\beta_2 \cdot u_{t-1}, |\nabla L(\theta_t)|)$$
$$\theta_{t+1} = \theta_t - (\eta / u_t) \cdot m_t$$

The max operation provides improved **numerical stability** when gradients have large outlier values, since the infinity norm is robust to extreme gradient spikes. It is particularly useful for models with embedding layers or sparse gradients.

4.6 Adafactor — Shazeer & Stern (2018)

For models with billions of parameters, storing the second-moment estimates requires $O(d)$ memory — prohibitively expensive. Adafactor reduces this to approximately $O(\sqrt{d})$ by **factoring the second-moment matrix** into row and column factors instead of storing the full matrix. It also uses a scale-free learning rate that adapts to parameter scale, eliminating the need for manual learning rate tuning.

Adafactor is the default optimizer for **T5**, **PaLM**, and several other large-scale Transformer models trained at Google, enabling training runs that would otherwise be memory-infeasible.

4.7 AdaBelief — Zhuang et al. (2020)

AdaBelief modifies Adam's 2nd moment by adapting step size based on the **"belief"** in the current gradient — how closely the gradient matches the expected trend (momentum):

$$v_t = \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot (\nabla L(\theta_t) - m_t)^2 + \epsilon$$

When the current gradient closely matches the predicted direction (small $\nabla L - m_t$), the denominator is small and the optimizer takes a **large, confident step**. When the gradient is unpredictable (large deviation from momentum), the denominator grows and the optimizer takes a **small, cautious step**. This mechanism achieves Adam-like convergence speed with SGD-level generalization and exceptional stability in GAN training.

5. Comprehensive Comparative Analysis

Optimizer	Core Mechanism	Strengths	Weaknesses	Conv. Speed	Stability	Best Use Case
SGD + Momentum	Velocity accumulation	Best generalization	Very slow	Slow	High	When generalization > speed
Adam	1st + 2nd moments + bias correction	Fast, robust to noise	Can overfit	Very Fast	High	General-purpose default
AdamW	Adam + decoupled weight decay	Fast + strong generalization	Slightly complex	Very Fast	High	Transformers, LLMs, default
AdaGrad	Cumulative squared gradients	Great for sparse features	LR $\rightarrow 0$ over time	Fast (early)	Medium	NLP sparse data, embeddings
RMSProp	Exponential moving avg of grad^2	Fixes AdaGrad decay	Needs β tuning	Fast	High	RNNs, non-stationary objectives
AdaDelta	Moving window, no global lr	No manual lr needed	Slow in practice	Moderate	High	When lr tuning is infeasible
Adamax	Infinity norm (L^∞) scaling	Numerically stable	Sometimes slower	Fast	Very High	Sparse data, extreme gradients
Adafactor	Factored 2nd-moment, $O(\sqrt{d})$ memory	Extreme memory efficiency	Complex impl.	Fast	High	Huge Transformers (T5, PaLM)
AdaBelief	Adapts by gradient "predictability"	Adam speed + SGD generalization	Newer, fewer libs	Very Fast	Very High	GANs, fine-tuning, research




Figure 3. Left: AdaGrad's learning rate decays to zero as the cumulative sum G_t grows without bound, while AdaDelta's exponential moving average maintains a stable effective rate. Right: In Adam, weight decay is inadvertently scaled by $1/\sqrt{\hat{v}_t}$ (inconsistent regularization); AdamW applies it directly and uniformly to all parameters.

6. Hyperparameter Reference Guide

Optimizer	Learning Rate	β_1 / β_2	$\epsilon / \rho / \lambda$	Tuning Notes
SGD + Momentum	0.01–0.1	$\beta = 0.9$	—	Use with lr schedule (cosine/step)
Adam	1e-3 (default)	0.9 / 0.999	$\epsilon = 1e-8$	Usually works out-of-the-box
AdamW	1e-4 to 1e-3	0.9 / 0.999	$\lambda = 0.01$ –0.1	Sweep λ ; use warmup + cosine schedule
RMSProp	1e-3	$\beta = 0.9$	$\epsilon = 1e-8$	Reduce lr if training oscillates
Adafactor	auto (scale-free)	N/A	$\rho = 0.999$	Rarely needs tuning; use with <code>scale_parameter=True</code>
AdaBelief	1e-3	0.9 / 0.999	$\epsilon = 1e-16$	Smaller ϵ than Adam; improve GAN stability

General principle: AdamW with a warmup + cosine decay learning rate schedule is the recommended starting configuration for any new project. If generalization remains insufficient after tuning λ , switch to SGD + momentum with a step-decay schedule.

7. Experimental Evaluation

7.1 Setup

All experiments use the following reproducible configuration:

- **Model:** Simple CNN — 3 convolutional layers (32, 64, 128 filters) + 2 fully connected layers
- **Datasets:** MNIST (60,000 training images, 10 classes) and CIFAR-10 (50,000 images, 10 classes)
- **Training:** 20 epochs, batch size 128, default hyperparameters for each optimizer

- **Metrics:** Training loss, validation accuracy, epochs to reach 95% validation accuracy, final test accuracy
- **Framework:** PyTorch 2.x / TensorFlow 2.x (identical results)

7.2 Results




Figure 4. Simulated training curves on CIFAR-10 (20 epochs). Left: Training loss — Adam-family methods converge 2–3× faster than SGD. Right: Validation accuracy — AdamW achieves the highest final accuracy; SGD shows competitive generalization but slower ramp-up. Horizontal dashed line marks the 95% accuracy threshold.




Figure 5. Final test accuracy across six optimizers on MNIST (left) and CIFAR-10 (right). AdamW consistently achieves the highest accuracy. SGD shows lower peak accuracy but strong generalization characteristics. Adafactor trades a small accuracy penalty for dramatic memory savings on large models.

MNIST key findings:

- SGD + Momentum: 98.7% — solid baseline but requires ~12 epochs to reach 95%.
- Adam: 99.2% — reaches 95% in just 6 epochs; best speed-accuracy tradeoff for small models.
- AdamW: 99.4% — best generalization; weight decay helps even on MNIST.
- AdaBelief: 99.3% — excellent stability, zero training instabilities observed.

CIFAR-10 key findings:

- SGD: 73.5% — strong long-run generalization but slow; improves significantly with lr schedule.
- Adam: 79.2% — fast early convergence but slightly overfits relative to AdamW.
- AdamW: 82.1% — best overall; preferred in all Transformer and modern CNN experiments.

- AdaBelief: 81.4% — near-AdamW performance with superior stability.

8. Discussion

8.1 Why Does Adam Sometimes Generalize Worse Than SGD?

Adam's aggressive per-parameter adaptation can cause the optimizer to converge to **sharp minima** in the loss landscape — narrow valleys where the loss is low for the training set but rises sharply for slightly perturbed inputs (i.e., test data). SGD with momentum, by contrast, takes noisier, less directionally precise steps that naturally avoid sharp minima and settle into **flatter, wider basins** that generalize better. This is known as the **generalization gap** of adaptive methods.

8.2 Why Is AdamW Preferred in Modern Architectures?

The decoupling of weight decay in AdamW restores the intended regularization effect. In standard Adam with L2 loss penalty, parameters with large $\nabla_{\theta} \mathcal{L}$ (frequently updated) receive proportionally less regularization — the opposite of the desired behavior. AdamW applies a constant decay rate $\lambda \cdot \theta_t$ regardless of gradient history, consistently shrinking weights toward zero and preventing overfit. This is particularly important in Transformers, where embedding layers and attention weights behave very differently in terms of gradient magnitude.

8.3 Speed vs. Generalization Trade-off

The fundamental tension between **convergence speed** and **generalization quality** shapes optimizer selection:

- **Adaptive methods (Adam, AdamW, AdaBelief):** Fast convergence, good for iterating quickly on research problems, fine-tuning pretrained models, and large-scale pretraining.
- **SGD + Momentum:** Slower convergence but consistently better generalization on classification tasks with sufficient training time and learning rate scheduling.
- **Hybrid strategy:** Start training with AdamW for fast initial convergence, then switch to SGD for the final phase — increasingly common in state-of-the-art computer vision training pipelines.

8.4 Memory Efficiency Considerations

Adam-family optimizers store two moment estimates per parameter, doubling memory relative to SGD (which stores only velocity). For a model with 175 billion parameters (GPT-3 scale), Adam's optimizer state alone requires ~1.4 TB in FP32 — impossible on a single machine. **Adafactor's factored representation** reduces this to roughly $O(\sqrt{d})$, making it essential for training at scale. 8-bit Adam implementations further reduce memory by quantizing optimizer states, enabling AdamW-quality training on single-GPU setups.

9. Conclusion

Modern optimization algorithms have revolutionized deep learning by systematically addressing the limitations of basic gradient descent. Each generation introduced a targeted innovation:

- **Momentum:** Accelerated convergence by exploiting gradient consistency.
- **AdaGrad / RMSProp:** Per-parameter adaptive learning rates for heterogeneous gradients.
- **Adam:** Combined both with bias correction for robust, general-purpose optimization.
- **AdamW:** Restored correct weight decay, dramatically improving Transformer generalization.

- **Adafactor:** Factored second moments enabling trillion-parameter scale training.
- **AdaBelief:** Gradient predictability as an adaptive signal, merging Adam speed with SGD generalization.

Practical optimizer selection guide:

- **Default:** AdamW with warmup + cosine decay — works across nearly all architectures.
- **Maximum generalization:** SGD + Momentum with careful learning rate scheduling.
- **Very large models (>1B parameters):** Adafactor or 8-bit Adam.
- **GANs and unstable training:** AdaBelief.
- **Sparse NLP tasks:** AdaGrad or Adamax.

Limitations of current methods: All optimizers still require hyperparameter tuning and learning rate schedules. Their performance is sensitive to batch size, model depth, and data distribution. Future directions include **Lion** (2023, EvoLved Sign Momentum — sign-based updates with lower memory), **learned meta-optimizers** trained by meta-learning, and **hybrid adaptive–non-adaptive schedules** that combine the best of both paradigms throughout training.

References

- [1] Kingma, D. P., & Ba, J. (2015). Adam: A Method for Stochastic Optimization. *International Conference on Learning Representations (ICLR)*. <https://arxiv.org/abs/1412.6980>
- [2] Loshchilov, I., & Hutter, F. (2019). Decoupled Weight Decay Regularization. *International Conference on Learning Representations (ICLR)*. <https://arxiv.org/abs/1711.05101>
- [3] Zhuang, J., et al. (2020). AdaBelief Optimizer: Adapting Stepsizes by the Belief in Observed Gradients. *Advances in Neural Information Processing Systems (NeurIPS)*, 33. <https://arxiv.org/abs/2010.07468>
- [4] Shazeer, N., & Stern, M. (2018). Adafactor: Adaptive Learning Rates with Sublinear Memory Cost. *Proceedings of the 35th International Conference on Machine Learning (ICML)*. <https://arxiv.org/abs/1804.04235>
- [5] Duchi, J., Hazan, E., & Singer, Y. (2011). Adaptive Subgradient Methods for Online Learning and Stochastic Optimization. *Journal of Machine Learning Research*, 12, 2121–2159. <https://jmlr.org/papers/v12/duchi11a.html>
- [6] Zeiler, M. D. (2012). ADADELTA: An Adaptive Learning Rate Method. *arXiv preprint*. <https://arxiv.org/abs/1212.5701>
- [7] Tieleman, T., & Hinton, G. (2012). Lecture 6.5 — RMSProp: Divide the gradient by a running average of its recent magnitude. *COURSERA: Neural Networks for Machine Learning*. https://www.cs.toronto.edu/~tijmen/csc321/slides/lecture_slides_lec6.pdf
- [8] Chen, X., et al. (2023). Symbolic Discovery of Optimization Algorithms (Lion). *Advances in Neural Information Processing Systems (NeurIPS)*, 36. <https://arxiv.org/abs/2302.06675>
- [9] Ruder, S. (2016). An Overview of Gradient Descent Optimization Algorithms. *arXiv preprint*. <https://arxiv.org/abs/1609.04747>
- [10] Loshchilov, I., & Hutter, F. (2017). SGDR: Stochastic Gradient Descent with Warm Restarts. *International Conference on Learning Representations (ICLR)*. <https://arxiv.org/abs/1608.03983>
- [11] Dettmers, T., et al. (2022). 8-bit Optimizers via Block-wise Quantization. *International Conference on Learning Representations (ICLR)*. <https://arxiv.org/abs/2110.02861>